

libdispatch

- Grand Central Dispatch
- Asynchronous & concurrent programming model
- From apple 
- <http://libdispatch.macosforge.org/>

Based on Queues

- split tasks to **blocks** and send them to different queues.
- A **block** is scheduled in its target queue.
- Notification when a group of **blocks** finish executing.
- Queue types: Global Concurrent Queues, Main Queue, Private Serial Queues

Global Concurrent Queues

- `q = dispatch_get_global_queue(
DISPATCH_QUEUE_PRIORITY_DEFAULT,
NULL /* reserved for future use */);`
- Execute function `complex_calculation` 100 times:
 - `dispatch_apply_f(100, q,
user_data, complex_calculation);`
 - `complex_calculation(user_data, i); /* i ∈ [0, 100) */`
 - more than one `complex_calculation` run parallelly

Main Queue

- Is a serial queue (back up by one thread)
- `q_main = dispatch_get_main_queue();`
- Is a global queue
- To integrate with Apple's Cocoa framework

Private Serial Queues

- `q_sum = dispatch_queue_create("com.example.sum", NULL);`
- **Serialize access to shared data structures:**

```
#define COUNT 128
```

```
double sum = 0;
```

```
void calc_func(void *data, size_t i) {
```

```
    double x = complex_calculation(i);
```

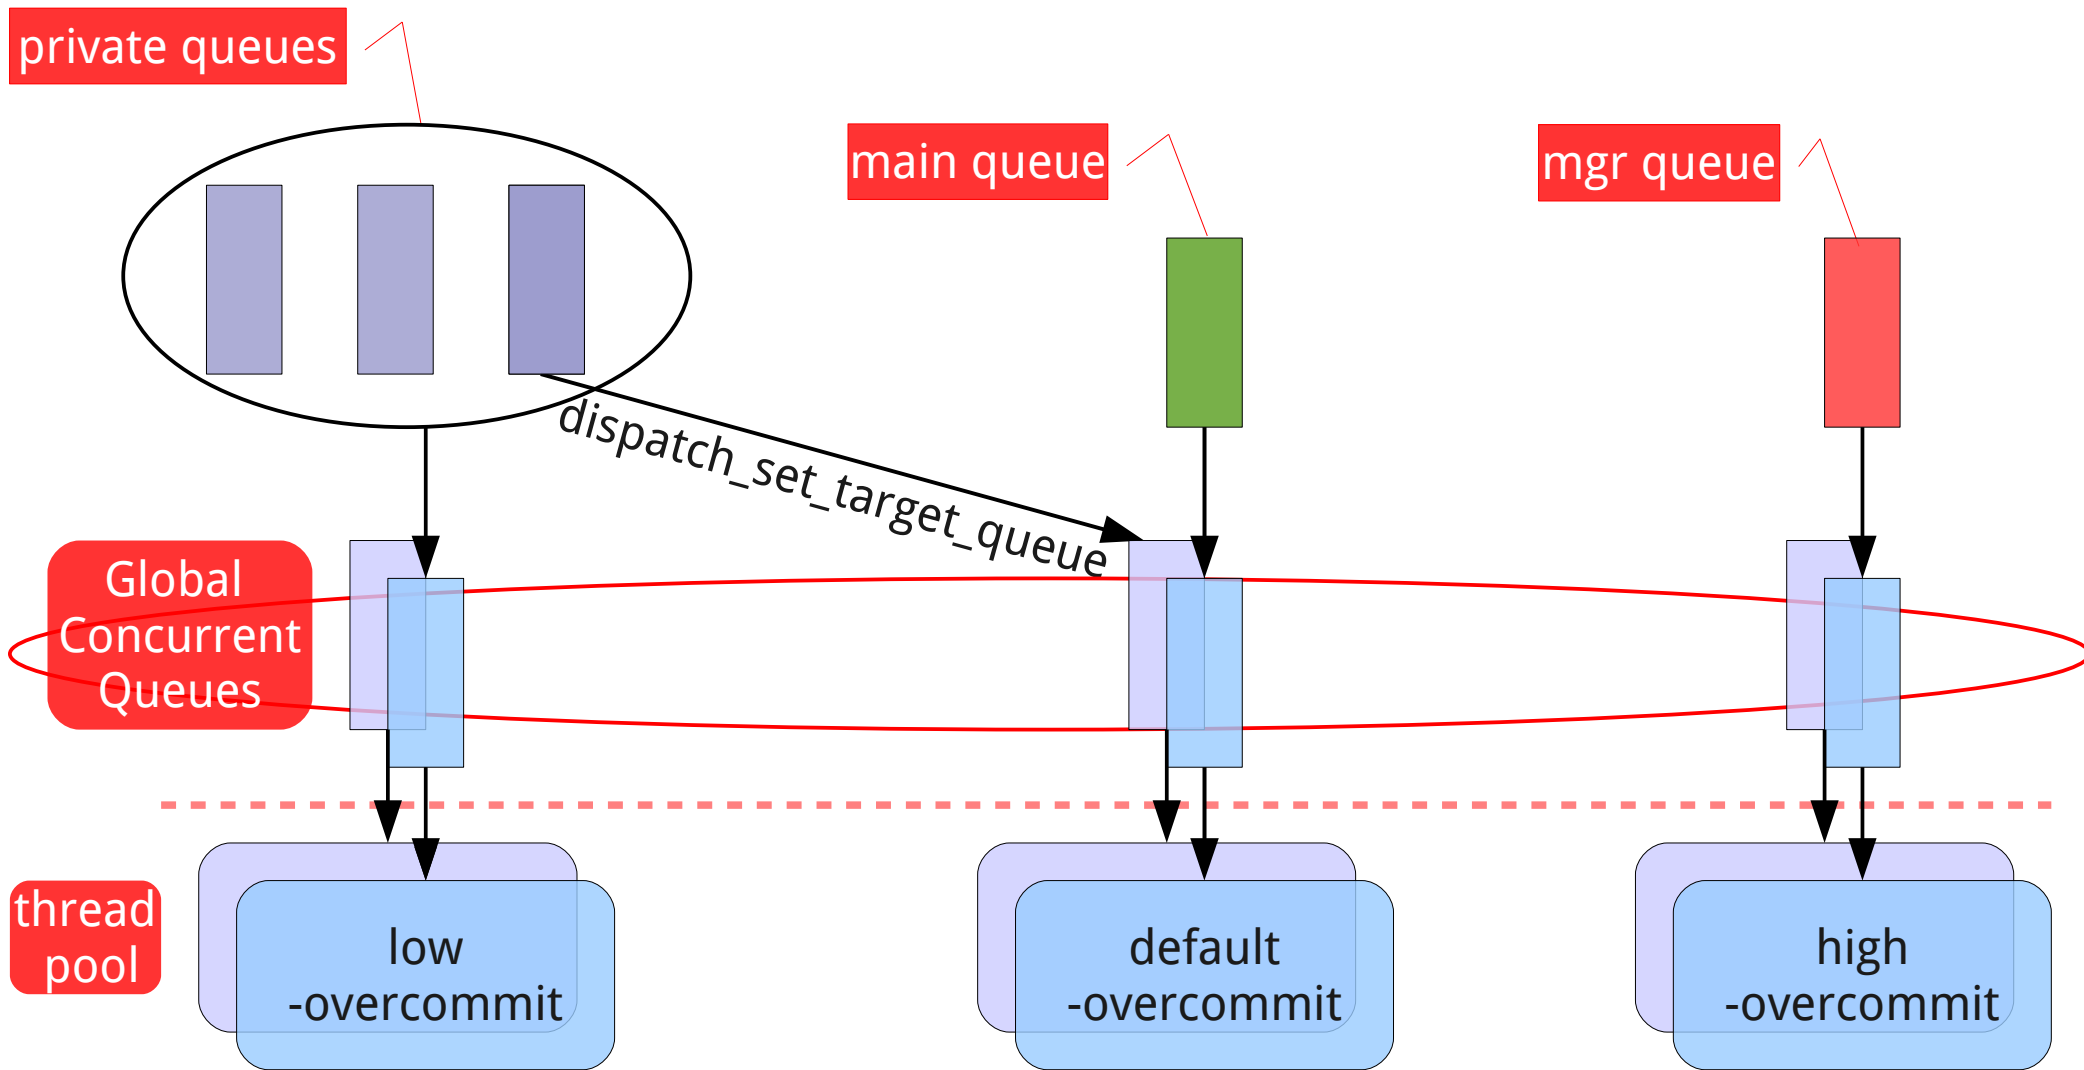
```
    double *sum = (double *)data;
```

```
    dispatch_async(q_sum, ^{ *sum += x});
```

```
}
```

```
dispatch_apply_f(COUNT, q_default, &sum, calc_func);
```

Relations between queues



Main classes and inheritance

dispatch_object_s

`const void *do_vtable;`
`struct x *volatile do_next;`

`unsigned int do_ref_cnt;`
`unsigned int do_xref_cnt;`
`unsigned int do_suspend_cnt;`
`struct dispatch_queue_s *do_targetq;`
`void *do_ctxt;`
`dispatch_function_t do_finalizer;`

dispatch_continuation_s

→ dispatch_queue_s

└─ dispatch_source_s

→ dispatch_queue_attr_s

→ dispatch_source_attr_s

→ dispatch_semaphore_s = dispatch_group_s

dispatch_queue_s

- Contain a list of DO(dispatch_object_s)

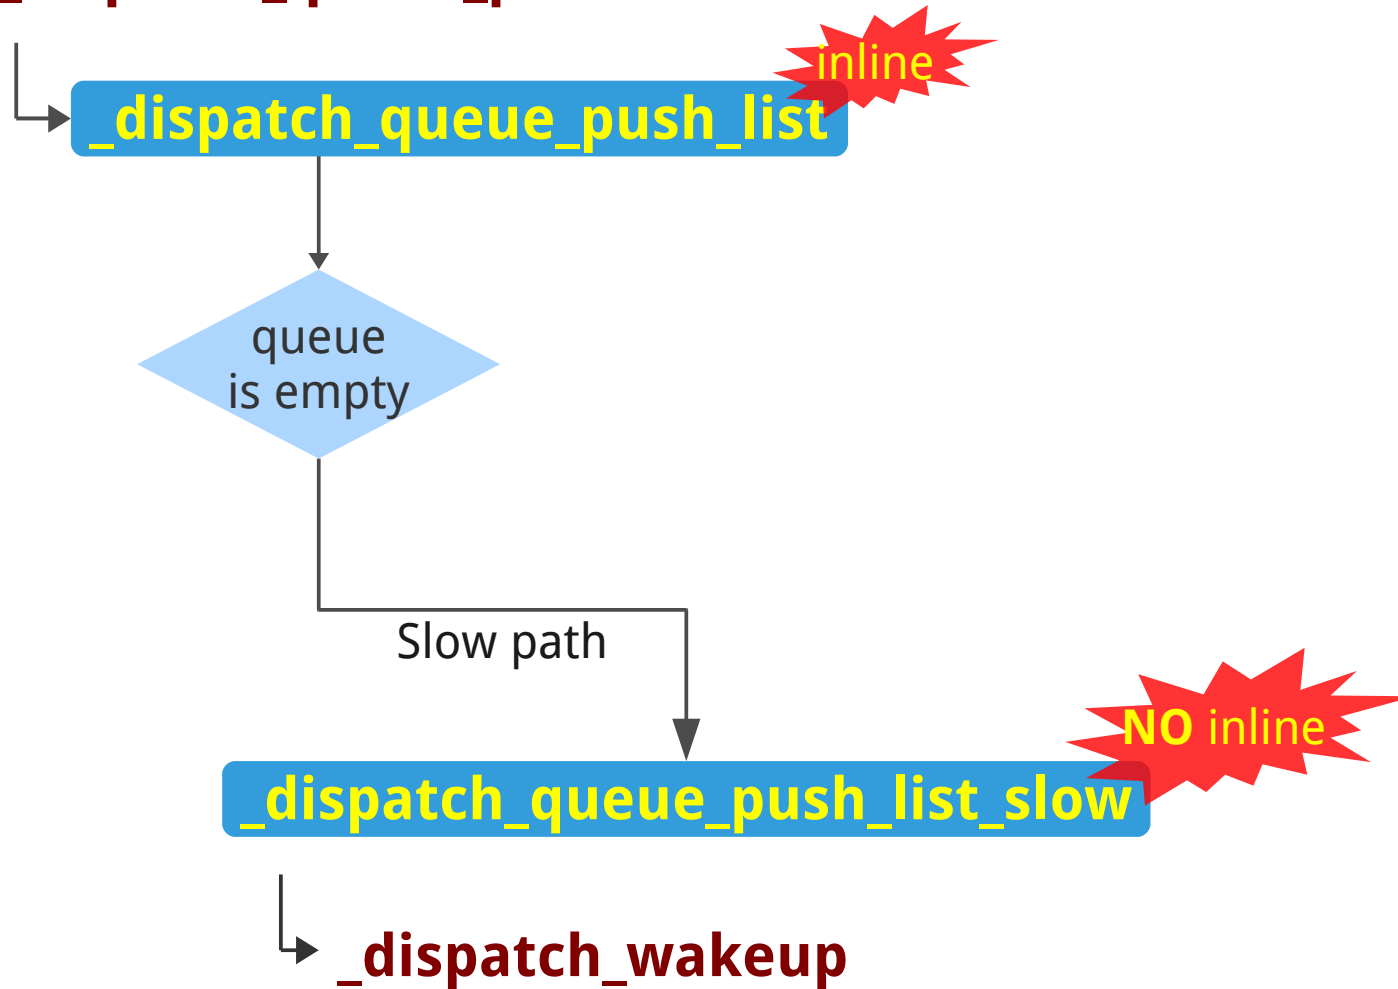
struct dispatch_object_s *
dq_items_head → DO → DO → DO → NULL

struct dispatch_object_s *
volatile dq_items_tail

- Num of Running DO: uint32_t dq_running;
- Width of concurrency: uint32_t dq_width;

Enqueue

_dispatch_queue_push



Deque

- `_dispatch_queue_concurrent_drain_one`
 - Get and return a DO concurrently
- `_dispatch_queue_drain`
 - Get and process all DOs in the queue
 - Lock the queue before calling:
`_dispatch_queue_trylock(dq)`

How a **block** be executed?

1. wrap a **block** to *dispatch_continuation_s*
2. **_dispatch_queue_push** to its target queue → **_dispatch_wakeup** the target queue if empty
3. **_dispatch_wakeup** do the following:
 - If SUSPENDED, return NULL
 - Run vtable->do_probe, if return false and the queue is empty, return NULL
 - **_dispatch_trylock** (object lock), if lock fail, return NULL
 - **_dispatch_queue_push**(dou.do->**do_targetq**, dou._do);
4. Finally **_dispatch_queue_push** to a root queue (i.e. Global Concurrent Queue, do_targetq == NULL)

Send to thread pool

_dispatch_wakeup(*root queue*)




Send to thread pool

_dispatch_wakeup(*root queue*)

└─ vtable->do_probe -----> **_dispatch_queue_wakeup_global**

int
pthread_workqueue_additem_np(
pthread_workqueue_t workq,
void *(*workitem_func)(void *), void * workitem_arg,
pthread_workitem_handle_t * itemhandlep, unsigned int *gencountp)



Send to thread pool

_dispatch_wakeup(*root queue*)

└─ vtable->do_probe -----> **_dispatch_queue_wakeup_global**



int
pthread_workqueue_additem_np(
pthread_workqueue_t workq,
void *(***workitem_func**)(void *), void * workitem_arg,
pthread_workitem_handle_t * itemhandlep, unsigned int *gencountp)

_dispatch_worker_thread2

└─ while ((item = fastpath(**_dispatch_queue_concurrent_drain_one**(dq))))
 _dispatch_continuation_pop(item);

Executing

- `_dispatch_continuation_pop`
 - Is a "dispatch_continuation_s" ?
 - Process flag: DISPATCH_OBJ_ASYNC_BIT
 - Process flag: DISPATCH_OBJ_GROUP_BIT
 - `dc->dc_func(dc->dc_ctxt)`
 - Or is a "dispatch_queue_s"?
 - Run `_dispatch_queue_invoke`
 1. Check SUSPEND state and try to acquire *queue lock*
 2. `_dispatch_queue_drain`
 3. Release *queue lock*
 4. Release *object lock* (locked in `_dispatch_wakeup`)

When wake up queues?

- push to an empty queue
- dq_running is 0
- _dispatch_queue_wakeup_global in
_dispatch_queue_concurrent_drain_one (fork
more working threads)

Implementation of thread pool

- Use Darwin's extension to POSIX threads
 - Create thread pool: `pthread_workqueue_create_np`
 - Adjust pool size by the overall load on the system
 - Add a job: `pthread_workqueue_additem_np`
- Built-in lightweight implementation
 - Pool size: `dgq_thread_pool_size`
 - Worker function: `_dispatch_worker_thread`
 - When all jobs complete, working thread will sleep on a signal several seconds, unless be waken up or quit on timeout

Other implementation technique

- Two reference counts
 - Internal reference count (`do_ref_cnt`)
 - External reference count (`do_xref_cnt`) – Better error detection for client code
- An simple but efficient memory allocation cache
 - Only cache `dispatch_continuation_t`
 - Per-thread, single link
 - Only flush cache on some points, usually when a working thread finishes all jobs
- fastpath, slowpath

Port to Linux

- By Mark Heily
- <http://packages.debian.org/squeeze/libdispatch0>
- Related libraries:
 1. [libkqueue](#) (implement kevent on top of epoll, inotify, signalfd and timerfd)
 2. [libpthread_workqueue](#) (implement pthread_workqueue in userspace)

END